

## 1. Cache principles

- a. Cache holds a much smaller subset of RAM's contents
  - i. Bring in values from RAM while executing a program
  - ii. If we need a value, look in cache first before going to memory
  - iii. Benefit from having cache because it is faster than RAM, and closer to the CPU
  - iv. However, cache is smaller than RAM, and costs more per bit than RAM
- b. Terminology
  - i. Memory *block* – unit of main memory stored in a cache line
  - ii. Cache *line* – basic unit of a cache
    - 1. Contains a block of memory
    - 2. Also contains a tag, and control bits to determine when line should be replaced
  - iii. Cache *tag* – number stored with a line
    - 1. Along with the line's position in the cache, determines the address of the block from main memory
  - iv. Cache *hit* – data that the CPU asks for is in the cache
  - v. Cache *miss* – data that the CPU asks for is *not* in the cache
  - vi. *Eviction* – removing memory from the cache and replacing with new memory
  - vii. *Dirty* – line contains updated data that differs from the corresponding main memory
    - 1. Before evicting the line from cache, must copy back to main memory
    - 2. Otherwise, would lose changes to that data

## 2. Cache design

- a. Cache design
  - i. Multiple-level caches
    - 1. See these in many modern CPUs
    - 2. L1 cache is the smallest and closest to the CPU
    - 3. L2 cache is larger than the L1, but is further away from the CPU
      - a. Takes more time to access the L2 than the L1
    - 4. Same applies for L3
      - a. Even larger than L2, but even slower
  - ii. Unified or split
    - 1. Split - have separate caches for data or instructions
    - 2. Unified – combine the two
- b. How to address the cache
  - i. Cache addresses can be *physical*, using actual memory addresses
    - 1. Slower, but avoids aliasing problem below
  - ii. Can also be *logical*, using the same virtual addresses that the program uses
    - 1. Memory management unit (MMU) converts virtual to physical addresses
    - 2. Faster, doesn't need to wait for MMU
    - 3. Encounters the *aliasing* problem
      - a. Same virtual address can be used in multiple user processes
      - b. Cache would say there's a hit
      - c. However, that data doesn't correspond to process currently running
      - d. Need to flush cache every time we switch processes
      - e. Incur a performance penalty every time we flush cache
- c. Cache size
  - i. Larger cache means slower cache
  - ii. There is no optimum cache size
    - 1. Exact specifications will always depend on machine being run